

AKS Networking Plugin

Kubenet vs Azure CNI



Houssem Dellai



Network Plugin options in AKS

Kubenet (Basic Networking)

Default mode for AKS

Pod CIDR

Internal network

Azure CNI (Advanced Networking)

Dynamic IP allocation

Overlay mode

Cilium

Bring Your Own (BYO)

Calico CNI

Network Plugin options in AKS

Kubenet (Basic Networking)

Default mode for AKS

Pod CIDR

Internal network

Azure CNI (Advanced Networking)

Dynamic IP allocation

Overlay mode

Cilium

Bring Your Own (BYO)

Calico CNI

Create Kubernetes cluster ...

Basics Node pools Access Networking Integrations Advanced Tags Review + create

You can choose between two networking options: 'Kubenet' or 'Azure CNI'.

- The **kubenet** networking plug-in creates a new VNet for your cluster using default values.
- The **Azure CNI** networking plug-in allows clusters to use a new or existing VNet with customizable addresses. Application pods are connected directly to the VNet, which allows for native integration with VNet features.

[Learn more about networking in Azure Kubernetes Service](#)

Network configuration ⓘ

☒ Kubenet

☐ Azure CNI

DNS name prefix * ⓘ

aks-dns ✓

Traffic routing


Load balancer ⓘ

Standard

Review + create

< Previous

Next : Integrations >

 Give feedback

AKS Network Plugin: **Kubenet**

Nodes get an IP address from the Azure virtual network subnet.

Pods receive an IP address from a logically different address space to the Azure virtual network subnet of the nodes.

Network address translation (**NAT**) is then configured so that the pods can reach resources on the VNET.








The source IP address of the traffic is NAT'd to the node's primary IP address.

This reduces the number of IP addresses that you need to reserve in your network space for pods to use.

Create an AKS cluster with Kubenet

```
az group create --name rg-aks-cni --location westeurope
```

```
az aks create -g rg-aks-cni -n aks-cni --network-plugin kubenet
```

<input type="checkbox"/> Name ↑↓	Type ↑↓
<input type="checkbox"/>  45508639-9b7d-468c-9eec-e0602fab461f	Public IP address
<input type="checkbox"/>  aks-agentpool-21301641-nsg	Network security group
<input type="checkbox"/>  aks-agentpool-21301641-routetable	Route table
<input type="checkbox"/>  aks-kubenet-agentpool	Managed Identity
<input type="checkbox"/>  aks-nodepool1-80782964-vmss	Virtual machine scale set
<input type="checkbox"/>  aks-vnet-21301641	Virtual network
<input type="checkbox"/>  kubernetes	Load balancer

Kubenet: Nodes IPs are from Subnet

```
$ kubectl get nodes -o wide
```

```
NAME
```

```
aks-nodepool1-10626751-vmss000000
```

```
aks-nodepool1-10626751-vmss000001
```

```
aks-nodepool1-10626751-vmss000002
```

```
$
```

```
INTERNAL-IP
```

```
10.224.0.5
```

```
10.224.0.6
```

```
10.224.0.4
```



aks-vnet-21301641 | Subnets



Virtual network

Search



+ Subnet

+ Gateway subnet

Settings

Address space

Connected devices

Subnets

Search subnets

Name ↑↓

IPv4 ↑↓

aks-subnet

10.224.0.0/16

Kubernetes: Pods IPs are from Pod CIDR (except few system Pods)

```
$ kubectl get nodes -o wide
NAME
aks-nodepool1-10626751-vmss000000
aks-nodepool1-10626751-vmss000001
aks-nodepool1-10626751-vmss000002
$
$ kubectl get pods -A -o wide
NAMESPACE      NAME
kube-system    azure-ip-masq-agent-4rh2p
kube-system    azure-ip-masq-agent-5w5r8
kube-system    azure-ip-masq-agent-sv8m4
kube-system    cloud-node-manager-7sblb
kube-system    cloud-node-manager-8fstz
kube-system    cloud-node-manager-jg826
kube-system    coredns-59b6bf8b4f-66hxq
kube-system    coredns-59b6bf8b4f-t8wnb
kube-system    coredns-autoscaler-5655d66f64-2mvl8
kube-system    csi-azuredisk-node-fxbnk
kube-system    csi-azuredisk-node-mcvzp
kube-system    csi-azuredisk-node-mprpw
kube-system    csi-azurefile-node-78npj
kube-system    csi-azurefile-node-cf6bl
kube-system    csi-azurefile-node-p9xbs
kube-system    konnectivity-agent-7cc8ddb5bc-flqkk
kube-system    konnectivity-agent-7cc8ddb5bc-wzfmt
kube-system    kube-proxy-6768z
kube-system    kube-proxy-hs997
kube-system    kube-proxy-k7kjk
kube-system    metrics-server-5f8d84558d-d5qhg
kube-system    metrics-server-5f8d84558d-wvjzf
$
```

INTERNAL-IP
10.224.0.5
10.224.0.6
10.224.0.4

IP
10.224.0.5
10.224.0.4
10.224.0.6
10.224.0.4
10.224.0.5
10.224.0.6
10.244.0.8
10.244.0.6
10.244.0.7
10.224.0.6
10.224.0.4
10.224.0.5
10.224.0.6
10.224.0.5
10.224.0.4
10.244.2.2
10.244.1.2
10.224.0.6
10.224.0.4
10.224.0.5
10.244.0.10
10.244.0.9

aks-vnet-21301641 | Subnets

Virtual network

Search

Subnet Gateway subnet

Settings

Address space

Connected devices

Subnets

Search subnets

Name	IPv4
aks-subnet	10.224.0.0/16

Networking

API server address	aks-kubene-rg-aks-kubenet-82f6d7-2zhbp6wv.hcp.westeurope.azmk8s.io
Network type (plugin)	Kubenet
Pod CIDR	10.244.0.0/16
Service CIDR	10.0.0.0/16
DNS service IP	10.0.0.10
Docker bridge CIDR	172.17.0.1/16


Kubernetes: Pods IPs are from Pod CIDR

```
$ kubectl get nodes -o wide
NAME
aks-nodepool1-10626751-vmss000000
aks-nodepool1-10626751-vmss000001
aks-nodepool1-10626751-vmss000002
$
$
$ kubectl create deployment nginx --image=nginx --replicas=10
deployment.apps/nginx created
$
$ kubectl get pods -o wide
NAME          READY   STATUS    RESTARTS   AGE   IP
nginx-8f458dc5b-2qp4n   1/1     Running   0           25s   10.244.2.5
nginx-8f458dc5b-6clrc   1/1     Running   0           26s   10.244.0.11
nginx-8f458dc5b-c9zrc   1/1     Running   0           25s   10.244.1.4
nginx-8f458dc5b-f7fdh   1/1     Running   0           25s   10.244.2.3
nginx-8f458dc5b-jgm7s   1/1     Running   0           25s   10.244.1.5
nginx-8f458dc5b-mw5gg   1/1     Running   0           25s   10.244.0.13
nginx-8f458dc5b-p4mn5   1/1     Running   0           25s   10.244.1.6
nginx-8f458dc5b-pnxwg   1/1     Running   0           26s   10.244.2.4
nginx-8f458dc5b-rckqv   1/1     Running   0           25s   10.244.0.12
nginx-8f458dc5b-tlqbp   1/1     Running   0           26s   10.244.1.3
$
```


INTERNAL-IP
10.224.0.5
10.224.0.6
10.224.0.4

```
$
$ kubectl get node -o jsonpath='{.items[*].spec.podCIDR}'
10.244.2.0/24 10.244.1.0/24 10.244.0.0/24
$
```

Each Node owns /24 CIDR for its Pods

 **Node pool**


Max pods per node
110


 **aks-vnet-21301641 | Subnets** ☆ ...


Virtual network

<< + Subnet + Gateway subnet


Settings

 Address space

 Connected devices

 Subnets

Name ↑↓	IPv4 ↑↓
aks-subnet	10.224.0.0/16

 **Networking**

API server address	aks-kubene-rg-aks-kubenet-82f6d7-2zhbp6wv.hcp.westeurope.azmk8s.io
Network type (plugin)	Kubenet
Pod CIDR	10.244.0.0/16
Service CIDR	10.0.0.0/16
DNS service IP	10.0.0.10
Docker bridge CIDR	172.17.0.1/16

Kubernetes: Services IPs are from Service CIDR (like CNI)

```
$ kubectl get nodes -o wide
NAME
aks-nodepool1-10626751-vmss000000
aks-nodepool1-10626751-vmss000001
aks-nodepool1-10626751-vmss000002
$
$
$ kubectl create deployment nginx --image=nginx --replicas=10
deployment.apps/nginx created
$
$ kubectl get pods -o wide
NAME                                READY   STATUS    RESTARTS   AGE   IP
nginx-8f458dc5b-2qp4n               1/1    Running   0           25s   10.244.2.5
nginx-8f458dc5b-6clrc               1/1    Running   0           26s   10.244.0.11
nginx-8f458dc5b-c9zrc               1/1    Running   0           25s   10.244.1.4
nginx-8f458dc5b-f7fdh               1/1    Running   0           25s   10.244.2.3
nginx-8f458dc5b-jgm7s               1/1    Running   0           25s   10.244.1.5
nginx-8f458dc5b-mw5gg               1/1    Running   0           25s   10.244.0.13
nginx-8f458dc5b-p4mn5               1/1    Running   0           25s   10.244.1.6
nginx-8f458dc5b-pnxwg               1/1    Running   0           26s   10.244.2.4
nginx-8f458dc5b-rckqv               1/1    Running   0           25s   10.244.0.12
nginx-8f458dc5b-tlqbp               1/1    Running   0           26s   10.244.1.3
$
$
$ kubectl get svc -A
NAMESPACE   NAME           TYPE        CLUSTER-IP
default     kubernetes     ClusterIP   10.0.0.1
default     nginx          ClusterIP   10.0.24.234
kube-system kube-dns       ClusterIP   10.0.0.10
kube-system metrics-server ClusterIP   10.0.228.174
$
```

aks-vnet-21301641 | Subnets

Virtual network

Search

Subnet Gateway subnet

Settings

Address space

Connected devices

Subnets

Search subnets

Name	IPv4
aks-subnet	10.224.0.0/16

Networking

API server address	aks-kubene-rg-aks-kubenet-82f6d7-2zhbp6wv.hcp.westeurope.azmk8s.io
Network type (plugin)	Kubenet
Pod CIDR	10.244.0.0/16
Service CIDR	10.0.0.0/16
DNS service IP	10.0.0.10
Docker bridge CIDR	172.17.0.1/16


How **Kubenet** works ?

Pods can't communicate directly with each other.

Instead, **User Defined Routing (UDR)** and IP forwarding is used for connectivity between pods across nodes.

The source IP address of the traffic is **NAT'd** to the node's primary IP address.

Pods from different nodes communicates through **Node route table**.



aks-agentpool-21301641-routetable

Route table

Search

Move Delete Refresh Give feedback

Essentials

JSON View

Routes

Search routes

Name	Address prefix	Next hop ty...	Next hop IP address
aks-nodepool1-10626751-vmss000000_...	10.244.2.0/24	Virtual appliance	10.224.0.5
aks-nodepool1-10626751-vmss000001_...	10.244.1.0/24	Virtual appliance	10.224.0.6
aks-nodepool1-10626751-vmss000002_...	10.244.0.0/24	Virtual appliance	10.224.0.4

Subnets

Search subnets

Name	Address range	Virtual network	Security group
aks-subnet	10.224.0.0/16	aks-vnet-21301641	aks-agentpool-21301641-nsg

Kubenet: Limitations

Azure supports a maximum of 400 routes in a UDR, so you can't have an AKS cluster larger than 400 nodes.

An additional hop is required in the design of kubenet, which adds minor latency to pod communication.

Kubenet doesn't support:

- AKS Virtual Nodes
- Windows Nodepools
- Azure Network Policies

You can use Calico Network Policies, as they are supported with kubenet.

Recap: AKS with Kubenet

Nodes are assigned IP addresses from Subnet.

Pods are assigned IP addresses from Pod CIDR (internal network in Kubernetes).

Some system Pods are assigned the same IP address of the Node !
Because Pods can request that.

AKS Network Plugin: Azure CNI

Every pod gets an IP address from the subnet and can be accessed directly.

These IP addresses must be unique across the network space and must be planned in advance.

Each node has a configuration parameter for the maximum number of pods that it supports.

The equivalent number of IP addresses per node are then reserved up front for that node.







This approach requires more planning.

Might lead to IP address exhaustion or the need to rebuild clusters in a larger subnet as application demands grow.

Create an AKS cluster with Azure CNI

```
az group create --name rg-aks-cni --location westeurope
```

```
az aks create -g rg-aks-cni -n aks-cni --network-plugin azure
```

<input type="checkbox"/> Name ↑↓	Type ↑↓
<input type="checkbox"/>  aks-agentpool-35006646-nsg	Network security group
<input type="checkbox"/>  aks-cni-agentpool	Managed Identity
<input type="checkbox"/>  aks-nodepool1-39057157-vmss	Virtual machine scale set
<input type="checkbox"/>  aks-vnet-35006646	Virtual network
<input type="checkbox"/>  b698b7dc-5b59-4dd0-b071-aae6a6463da3	Public IP address
<input type="checkbox"/>  kubernetes	Load balancer

Azure CNI : Node IPs are from Subnet

```
$  
$ kubectl get nodes -o wide  
NAME  
aks-nodepool1-35527294-vmss000000  
aks-nodepool1-35527294-vmss000001  
aks-nodepool1-35527294-vmss000002  
$
```

INTERNAL-IP
10.224.0.62
10.224.0.33
10.224.0.4

Node pool
Max pods per node 30

<> aks-vnet-21301641 | Subnets ☆ ...
Virtual network

Search << + Subnet + Gateway subnet

Settings

- <> Address space
- Connected devices
- <> Subnets

Search subnets


Name ↑↓	IPv4 ↑↓
aks-subnet	10.224.0.0/16

Azure CNI : Pods IPs are from Subnet

```
$  
$ kubectl get nodes -o wide
```

```
NAME  
aks-nodepool1-35527294-vmss000000  
aks-nodepool1-35527294-vmss000001  
aks-nodepool1-35527294-vmss000002
```

```
INTERNAL-IP  
10.224.0.62  
10.224.0.33  
10.224.0.4
```

 Node pool

Max pods per node
30


```
$  
$ kubectl get pods -A -o wide
```

```
NAMESPACE    NAME  
kube-system   azure-ip-masq-agent-bw2gh  
kube-system   azure-ip-masq-agent-v2z8c  
kube-system   azure-ip-masq-agent-xgfkf  
kube-system   cloud-node-manager-95mp2  
kube-system   cloud-node-manager-mnxwx  
kube-system   cloud-node-manager-xjbwd  
kube-system   coredns-59b6bf8b4f-49xf2  
kube-system   coredns-59b6bf8b4f-bdnkj  
kube-system   coredns-autoscaler-5655d66f64-955b5  
kube-system   csi-azuredisk-node-478kj  
kube-system   csi-azuredisk-node-fzqfb  
kube-system   csi-azuredisk-node-m8clc  
kube-system   csi-azurefile-node-j27tj  
kube-system   csi-azurefile-node-nf5d9  
kube-system   csi-azurefile-node-w7dqk  
kube-system   konnectivity-agent-7d8d9bfc4b-jgpbg  
kube-system   konnectivity-agent-7d8d9bfc4b-twvvq  
kube-system   kube-proxy-bn8gj  
kube-system   kube-proxy-q4r6t  
kube-system   kube-proxy-v9mjp  
kube-system   metrics-server-5f8d84558d-csnwl  
kube-system   metrics-server-5f8d84558d-z9wbm
```

```
IP  
10.224.0.4  
10.224.0.62  
10.224.0.33  
10.224.0.62  
10.224.0.4  
10.224.0.33  
10.224.0.67  
10.224.0.5  
10.224.0.63  
10.224.0.33  
10.224.0.62  
10.224.0.4  
10.224.0.33  
10.224.0.4  
10.224.0.62  
10.224.0.40  
10.224.0.23  
10.224.0.4  
10.224.0.33  
10.224.0.62  
10.224.0.17  
10.224.0.10
```

 aks-vnet-21301641 | Subnets ☆ ...


Virtual network

 Search

+ Subnet


+ Gateway subnet

Settings

 Address space

 Connected devices

 Subnets

 Search subnets

Name ↑↓

IPv4 ↑↓

aks-subnet

10.224.0.0/16

Networking

API server address aks-cni-rg-aks-cni-82f6d7-8oos0rk9.hcp.westeurope.azmk8s.io

Network type (plugin) Azure CNI

Pod CIDR

-

Service CIDR

10.0.0.0/16

DNS service IP

10.0.0.10

Docker bridge CIDR

172.17.0.1/16

Azure CNI : Services IPs are from Service CIDR (like Kubenet)

```
$  
$ kubectl get nodes -o wide  
NAME  
aks-nodepool1-35527294-vmss000000  
aks-nodepool1-35527294-vmss000001  
aks-nodepool1-35527294-vmss000002
```

INTERNAL-IP
10.224.0.62
10.224.0.33
10.224.0.4

Node pool
Max pods per node 30

```
$  
$ kubectl get pods -A -o wide
```

NAMESPACE	NAME
kube-system	azure-ip-masq-agent-bw2gh
kube-system	azure-ip-masq-agent-v2z8c
kube-system	azure-ip-masq-agent-xgfmt
kube-system	cloud-node-manager-95mp2
kube-system	cloud-node-manager-mnxwx
kube-system	cloud-node-manager-xjbwd
kube-system	coredns-59b6bf8b4f-49xf2
kube-system	coredns-59b6bf8b4f-bdnkj
kube-system	coredns-autoscaler-5655d66f64-955b5
kube-system	csi-azuredisk-node-478kj
kube-system	csi-azuredisk-node-fzqfb
kube-system	csi-azuredisk-node-m8clc
kube-system	csi-azurefile-node-j27tj
kube-system	csi-azurefile-node-nf5d9
kube-system	csi-azurefile-node-w7dnk

IP
10.224.0.4
10.224.0.62
10.224.0.33
10.224.0.62
10.224.0.4
10.224.0.33
10.224.0.67
10.224.0.5
10.224.0.63
10.224.0.33
10.224.0.62
10.224.0.4
10.224.0.33
10.224.0.4
10.224.0.62

```
$  
$ kubectl get services -A
```

NAMESPACE	NAME	TYPE	CLUSTER-IP
default	kubernetes	ClusterIP	10.0.0.1
kube-system	kube-dns	ClusterIP	10.0.0.10
kube-system	metrics-server	ClusterIP	10.0.249.103

aks-vnet-21301641 | Subnets ☆ ...
Virtual network

Search << + Subnet + Gateway subnet

Settings

<> Address space

Connected devices

<> Subnets

Search subnets

Name ↑↓	IPv4 ↑↓
aks-subnet	10.224.0.0/16

Networking

API server address aks-cni-rg-aks-cni-82f6d7-8oos0rk9.hcp.westeurope.azmk8s.io

Network type (plugin)	Azure CNI
Pod CIDR	-
Service CIDR	10.0.0.0/16
DNS service IP	10.0.0.10
Docker bridge CIDR	172.17.0.1/16

IPs utilization in **Kubenet** vs **Azure CNI**

How many Nodes and Pods could be deployed into a Subnet with CIDR /24 ?

Subnet reserves the first three IPs for management operations.

Default maximum of 110 pods per node with kubenet.

Default maximum of 30 pods per node with Azure CNI.

Subnet CIDR /24	Kubenet	Azure CNI
#Nodes	251 nodes	8 nodes
#Pods	27,610 (251*110) pods	240 (8*30) pods

⇒ Problem of IP exhaustion in **Azure CNI**.

⇒ Could be resolved using **Azure CNI Overlay**.

Kubenet or Azure CNI ?

Use **kubenet** when:


- You have limited IP address space.
- Most of the pod communication is within the cluster.
- You don't need advanced AKS features such as virtual nodes or Azure Network Policy.

Use **Azure CNI** when:


- You have available IP address space.
- Most of the pod communication is to resources outside of the cluster.
- You don't want to manage user defined routes for pod connectivity.
- You need AKS advanced features such as virtual nodes or Azure Network Policy.

Kubenet vs Azure CNI

Kubenet



Azure CNI



Azure CNI Overlay overview

In overlay networking, only the Kubernetes cluster nodes are assigned IPs from a subnet.

Pods receive IPs from a private CIDR that is provided at the time of cluster creation.

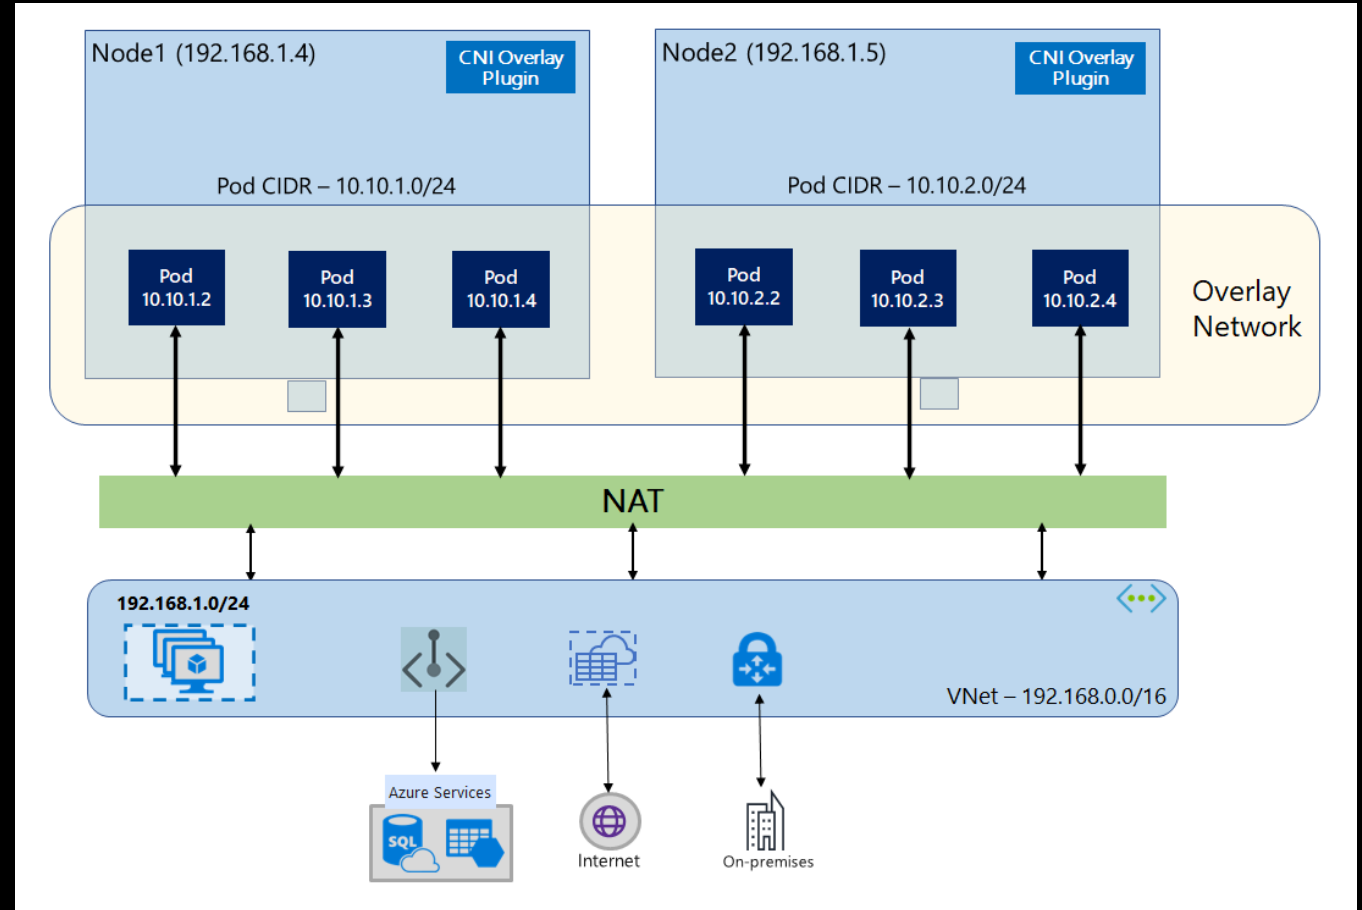
Each node is assigned:

- An IP address from Subnet
- /24 from Pod-CIDR to assign to its Pods.

There is no need to provision custom routes on the cluster subnet or use an encapsulation method to tunnel traffic between pods.

Promises connectivity **performance** between pods.







Endpoints outside the cluster can't connect to a pod directly.



Create cluster with Azure CNI Overlay

```
$ az group create -n rg-aks-cni-overlay -l westeurope
```


```
$ az aks create -n aks-cni-overlay -g rg-aks-cni-overlay `
  --network-plugin azure `
  --network-plugin-mode overlay `
  --pod-cidr 192.168.0.0/16
```

<input type="checkbox"/> Name ↑↓	Type ↑↓
<input type="checkbox"/>  aks-agentpool-51020892-nsg	Network security group
<input type="checkbox"/>  aks-cni-overlay-agentpool	Managed Identity
<input type="checkbox"/>  aks-nodepool1-14046306-vmss	Virtual machine scale set
<input type="checkbox"/>  aks-vnet-51020892	Virtual network
<input type="checkbox"/>  e1d2d241-1296-4c89-90ba-b300c5a5405d	Public IP address
<input type="checkbox"/>  kubernetes	Load balancer

Azure CNI Overlay: Node IPs are from Subnet

```
$  
$ kubectl get nodes -o wide  
NAME  
aks-nodepool1-14046306-vmss000000  
aks-nodepool1-14046306-vmss000001  
aks-nodepool1-14046306-vmss000002
```

INTERNAL-IP
10.224.0.4
10.224.0.5
10.224.0.6


 **Node pool**
Max pods per node
250

```
$  
$ kubectl get pods -A -o wide  
NAMESPACE      NAME  
kube-system     azure-cns-8p4pq  
kube-system     azure-cns-l76pf  
kube-system     azure-cns-mjz6p  
kube-system     azure-ip-masq-agent-9h2dw  
kube-system     azure-ip-masq-agent-ldwc8  
kube-system     azure-ip-masq-agent-lr64z  
kube-system     cloud-node-manager-d7p55  
kube-system     cloud-node-manager-fhkmr  
kube-system     cloud-node-manager-pcqdn  
kube-system     coredns-59b6bf8b4f-2nz4d  
kube-system     coredns-59b6bf8b4f-shk6w  
kube-system     coredns-autoscaler-5655d66f64-wwfdg  
kube-system     csi-azuredisk-node-6wpqt  
kube-system     csi-azuredisk-node-d992s  
kube-system     csi-azuredisk-node-fxgn2  
kube-system     csi-azurefile-node-74vcl  
kube-system     csi-azurefile-node-f7q72  
kube-system     csi-azurefile-node-zchd4  
kube-system     konnectivity-agent-79bc94d486-d4t9j  
kube-system     konnectivity-agent-79bc94d486-wb4rk  
kube-system     kube-proxy-p28kk  
kube-system     kube-proxy-p2k5c  
kube-system     kube-proxy-q6d4b  
kube-system     metrics-server-5f8d84558d-dt95p  
kube-system     metrics-server-5f8d84558d-jv46s
```

IP
10.224.0.6
10.224.0.4
10.224.0.5
10.224.0.5
10.224.0.6
10.224.0.4
10.224.0.6
10.224.0.5
10.224.0.5
10.224.0.4
192.168.1.33
192.168.0.165
192.168.1.142
10.224.0.5
10.224.0.4
10.224.0.6
10.224.0.6
10.224.0.4
10.224.0.5
192.168.1.53
192.168.1.62
10.224.0.6
10.224.0.4
10.224.0.5
192.168.0.233
192.168.0.151

aks-vnet-51020892 | Subnets


Virtual network

 Search

+ Subnet


+ Gateway subnet

Settings

 Address space

 Connected devices

 Subnets

 Search subnets

Name ↑↓

IPv4 ↑↓

aks-subnet

10.224.0.0/16



Networking

API server address aks-cni-ov-rg-aks-cni-overl-82f6d7-iznmpgpa.hcp.westeurope.azmk8s.io

Network type (plugin) Azure CNI

Pod CIDR -

Service CIDR 10.0.0.0/16


DNS service IP 10.0.0.10

Docker bridge CIDR 172.17.0.1/16

```
"networkProfile": {  
  "networkPlugin": "azure",  
  "networkPluginMode": "overlay",  
  "networkDataplane": "azure",  
  "loadBalancerSku": "Standard",  
  "loadBalancerProfile": { ...  
},  
  "podCidr": "192.168.0.0/16",  
  "serviceCidr": "10.0.0.0/16",  
  "dnsServiceIP": "10.0.0.10",  
  "dockerBridgeCidr": "172.17.0.1/16",  
}
```

Azure CNI Overlay: Pod IPs are from Pod-CIDR

```
$  
$ kubectl get nodes -o wide  
NAME                                INTERNAL-IP  
aks-nodepool1-14046306-vmss000000  10.224.0.4  
aks-nodepool1-14046306-vmss000001  10.224.0.5  
aks-nodepool1-14046306-vmss000002  10.224.0.6  
$
```


 Node pool



Max pods per node
250

```
$  
$ kubectl get pods -o wide  
NAME                                READY   IP  
nginx-8f458dc5b-4ddwt              1/1    192.168.1.143  
nginx-8f458dc5b-56qgw              1/1    192.168.2.41  
nginx-8f458dc5b-6n799              1/1    192.168.0.97  
nginx-8f458dc5b-d6dpx              1/1    192.168.1.174  
nginx-8f458dc5b-dcr5w              1/1    192.168.2.158  
nginx-8f458dc5b-fm7bj              1/1    192.168.2.190  
nginx-8f458dc5b-k26fp              1/1    192.168.0.220  
nginx-8f458dc5b-lm9wn              1/1    192.168.1.29  
nginx-8f458dc5b-mwfr2              1/1    192.168.0.171  
nginx-8f458dc5b-tvxgq              1/1    192.168.2.79  
$
```


aks-vnet-51020892 | Subnets

Virtual network

 Search


 Subnet  Gateway subnet

Settings

 Address space

 Connected devices

 Subnets

 Search subnets

Name ↑↓

IPv4 ↑↓

aks-subnet

10.224.0.0/16



Networking

API server address aks-cni-ov-rg-aks-cni-overl-82f6d7-
iznmpgpa.hcp.westeurope.azmk8s.io

Network type (plugin) Azure CNI

Pod CIDR -

Service CIDR 10.0.0.0/16

DNS service IP 10.0.0.10

Docker bridge CIDR 172.17.0.1/16

```
"networkProfile": {  
  "networkPlugin": "azure",  
  "networkPluginMode": "overlay",  
  "networkDataplane": "azure",  
  "loadBalancerSku": "Standard",  
  "loadBalancerProfile": { ...  
},  
  "podCidr": "192.168.0.0/16",  
  "serviceCidr": "10.0.0.0/16",  
  "dnsServiceIP": "10.0.0.10",  
  "dockerBridgeCidr": "172.17.0.1/16",  
}
```


Azure CNI : Service IPs are from Service-CIDR

```
$  
$ kubectl get nodes -o wide
```


NAME	INTERNAL-IP
aks-nodepool1-14046306-vmss000000	10.224.0.4
aks-nodepool1-14046306-vmss000001	10.224.0.5
aks-nodepool1-14046306-vmss000002	10.224.0.6



 **Node pool**

Max pods per node
250


aks-vnet-51020892 | Subnets

Virtual network

 Search


 Subnet  Gateway subnet

Settings

 Address space

 Connected devices

 Subnets

 Search subnets

Name ↑↓

IPv4 ↑↓

aks-subnet

10.224.0.0/16



Networking

API server address aks-cni-ov-rg-aks-cni-overl-82f6d7-
iznmpgpa.hcp.westeurope.azmk8s.io

Network type (plugin) Azure CNI

Pod CIDR

-

Service CIDR

10.0.0.0/16

DNS service IP

10.0.0.10

Docker bridge CIDR

172.17.0.1/16

```
"networkProfile": {  
  "networkPlugin": "azure",  
  "networkPluginMode": "overlay",  
  "networkDataplane": "azure",  
  "loadBalancerSku": "Standard",  
  "loadBalancerProfile": { ...  
},  
  "podCidr": "192.168.0.0/16",  
  "serviceCidr": "10.0.0.0/16",  
  "dnsServiceIP": "10.0.0.10",  
  "dockerBridgeCidr": "172.17.0.1/16",  
}
```

```
$  
$ kubectl get pods -o wide
```

NAME	READY	IP
nginx-8f458dc5b-4ddwt	1/1	192.168.1.143
nginx-8f458dc5b-56qgw	1/1	192.168.2.41
nginx-8f458dc5b-6n799	1/1	192.168.0.97
nginx-8f458dc5b-d6dpx	1/1	192.168.1.174
nginx-8f458dc5b-dcr5w	1/1	192.168.2.158
nginx-8f458dc5b-fm7bj	1/1	192.168.2.190
nginx-8f458dc5b-k26fp	1/1	192.168.0.220
nginx-8f458dc5b-lm9wn	1/1	192.168.1.29
nginx-8f458dc5b-mwfr2	1/1	192.168.0.171
nginx-8f458dc5b-tvxgq	1/1	192.168.2.79

```
$  
$ kubectl get services -A
```

NAMESPACE	NAME	TYPE	CLUSTER-IP
default	kubernetes	ClusterIP	10.0.0.1
kube-system	kube-dns	ClusterIP	10.0.0.10
kube-system	metrics-server	ClusterIP	10.0.67.197

Kubenet vs Azure CNI Overlay

Area	Azure CNI Overlay	Kubenet
Cluster scale	1000 nodes and 250 pods/node	400 nodes and 250 pods/node
Network configuration	Simple - no additional configuration required for pod networking	Complex - requires route tables and UDRs on cluster subnet for pod networking
Pod connectivity performance	Performance on par with VMs in a VNet	Additional hop adds minor latency
Kubernetes Network Policies	Azure Network Policies, Calico, Cilium	Calico
OS platforms supported	Linux and Windows Server 2022	Linux only

Limitations of Azure CNI Overlay

Azure CNI Overlay has the following limitations:

- You can't use Application Gateway as an Ingress Controller (AGIC) for an overlay cluster.
- Windows Server 2019 node pools are not supported for overlay.

Bring your own CNI

Cilium (OSS & Enterprise)

Calico CNI

Canal

Flannel

Weave

CIDR Subnet, Pod, Service & Docker Bridge overlapping

Subnet, Pod, Service & Docker Bridge should have different CIDR ranges. **No overlap.**
Pod, Service and Docker Bridge CIDRs could be the same for different clusters.

These IP address ranges should be an address space that isn't in use elsewhere in:

- Your network environment
 - including any on-premises network ranges connected, or plan to connect.
- Your Azure virtual networks using Express Route or a Site-to-Site VPN connection.

AKS clusters may not use:

- 169.254.0.0/16
- 172.30.0.0/16
- 172.31.0.0/16
- 192.0.2.0/24

for the :

- Service-CIDR
- Pod-CIDR
- Cluster virtual network

Can't be updated after cluster creation.

```
az aks create \  
  --resource-group myResourceGroup \  
  --name myAKSCluster \  
  --network-plugin kubenet \  
  --service-cidr 10.0.0.0/16 \  
  --dns-service-ip 10.0.0.10 \  
  --pod-cidr 10.244.0.0/16 \  
  --docker-bridge-address 172.17.0.1/16 \  
  --vnet-subnet-id $SUBNET_ID
```